

# Customer Churn Prediction

**Objective:** To develop a predictive model using the provided churn dataset that accurately identifies customers likely to exit

**Members:** Tyler, Humza, Angelina, Christine, Efaz, Nathan

In [ ]:

```
# General Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import missingno as ms
%matplotlib inline

# Machine Learning Libraries
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score, precision_score, recall_score, classification_report
```

## Part 1: Data Understanding

### Exploratory Data Analysis (EDA)

Let's take a look at the dataset and view only the first 10 rows

In [ ]:

```
churn = pd.read_csv('Churn.csv')
print("Size of Churn Set: ", churn.shape)
churn.head(10)
```

Size of Churn Set: (10000, 14)

Out[ ]:

|   | RowNumber | CustomerId | Surname  | CreditScore | Geography | Gender | Age | Tenure | Balance   | NumOfProducts | HasCrCard |
|---|-----------|------------|----------|-------------|-----------|--------|-----|--------|-----------|---------------|-----------|
| 0 | 1         | 15634602   | Hargrave | 619         | France    | Female | 42  | 2      | 0.00      | 1             |           |
| 1 | 2         | 15647311   | Hill     | 608         | Spain     | Female | 41  | 1      | 83807.86  | 1             |           |
| 2 | 3         | 15619304   | Onio     | 502         | France    | Female | 42  | 8      | 159660.80 | 3             |           |
| 3 | 4         | 15701354   | Boni     | 699         | France    | Female | 39  | 1      | 0.00      | 2             |           |
| 4 | 5         | 15737888   | Mitchell | 850         | Spain     | Female | 43  | 2      | 125510.82 | 1             |           |
| 5 | 6         | 15574012   | Chu      | 645         | Spain     | Male   | 44  | 8      | 113755.78 | 2             |           |
| 6 | 7         | 15592531   | Bartlett | 822         | France    | Male   | 50  | 7      | 0.00      | 2             |           |
| 7 | 8         | 15656148   | Obinna   | 376         | Germany   | Female | 29  | 4      | 115046.74 | 4             |           |
| 8 | 9         | 15792365   | He       | 501         | France    | Male   | 44  | 4      | 142051.07 | 2             |           |
| 9 | 10        | 15592389   | H?       | 684         | France    | Male   | 27  | 2      | 134603.88 | 1             |           |

Our target variable seems to be the Exited column which has responses containing either 1 for yes and 0 for no. Therefore we are dealing with a Binary Classification problem

It seems the dataset has 14 columns, with two of them likely not being two important: RowNumber and CustomerId, so let's drop the two columns

In [ ]:

```
# Dropping the first two columns as it doesn't offer much outside of being labels for each customer
churn = churn.drop(columns=['RowNumber', 'CustomerId', 'Surname'])
churn.head(5)
```

Out[ ]:

|   | CreditScore | Geography | Gender | Age | Tenure | Balance   | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary |
|---|-------------|-----------|--------|-----|--------|-----------|---------------|-----------|----------------|-----------------|
| 0 | 619         | France    | Female | 42  | 2      | 0.00      | 1             | 1         | 1              | 101348.88       |
| 1 | 608         | Spain     | Female | 41  | 1      | 83807.86  | 1             | 0         | 1              | 112542.58       |
| 2 | 502         | France    | Female | 42  | 8      | 159660.80 | 3             | 1         | 0              | 113931.57       |
| 3 | 699         | France    | Female | 39  | 1      | 0.00      | 2             | 0         | 0              | 93826.63        |
| 4 | 850         | Spain     | Female | 43  | 2      | 125510.82 | 1             | 1         | 1              | 79084.10        |

Now let's take a look at the data types for each of the variables in the dataframe

In [ ]:

```
data_types = churn.dtypes

print("Data Types and Additional Information:")
for column in churn.columns:
    print(f'Column: {column}')
    print(f' - Data Type: {data_types[column]}')
    print(f' - Number of Unique Values: {churn[column].nunique()}')
    print(f' - Sample Values: {churn[column].dropna().unique()[:5]}')
```

```
Data Types and Additional Information:
Column: CreditScore
- Data Type: int64
- Number of Unique Values: 460
- Sample Values: [619 608 502 699 850]
Column: Geography
- Data Type: object
- Number of Unique Values: 3
- Sample Values: ['France' 'Spain' 'Germany']
Column: Gender
- Data Type: object
- Number of Unique Values: 2
- Sample Values: ['Female' 'Male']
Column: Age
- Data Type: int64
- Number of Unique Values: 70
- Sample Values: [42 41 39 43 44]
Column: Tenure
- Data Type: int64
- Number of Unique Values: 11
- Sample Values: [2 1 8 7 4]
Column: Balance
- Data Type: float64
- Number of Unique Values: 6382
- Sample Values: [    0.    83807.86 159660.8 125510.82 113755.78]
Column: NumOfProducts
- Data Type: int64
- Number of Unique Values: 4
- Sample Values: [1 3 2 4]
Column: HasCrCard
```

- Data Type: int64
- Number of Unique Values: 2
- Sample Values: [1 0]

Column: IsActiveMember

- Data Type: int64
- Number of Unique Values: 2
- Sample Values: [1 0]

Column: EstimatedSalary

- Data Type: float64
- Number of Unique Values: 9999
- Sample Values: [101348.88 112542.58 113931.57 93826.63 79084.1 ]

Column: Exited

- Data Type: int64
- Number of Unique Values: 2
- Sample Values: [1 0]

**Numerical Variables: CreditScore, Age, Tenure, Balance, NumOfProducts, HasCrCard,IsActiveMember,EstimatedSalary,Exited** **Categorical Variables: Surname, Geography, Gender**

In [ ]:

```
churn.describe()
```

Out[ ]:

|       | CreditScore  | Age          | Tenure       | Balance       | NumOfProducts | HasCrCard    | IsActiveMember | EstimatedSal |
|-------|--------------|--------------|--------------|---------------|---------------|--------------|----------------|--------------|
| count | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000  | 10000.000000  | 10000.000000 | 10000.000000   | 10000.000000 |
| mean  | 650.528800   | 38.921800    | 5.012800     | 76485.889288  | 1.530200      | 0.70550      | 0.515100       | 100090.2390  |
| std   | 96.653299    | 10.487806    | 2.892174     | 62397.405202  | 0.581654      | 0.45584      | 0.499797       | 57510.4920   |
| min   | 350.000000   | 18.000000    | 0.000000     | 0.000000      | 1.000000      | 0.00000      | 0.000000       | 11.5800      |
| 25%   | 584.000000   | 32.000000    | 3.000000     | 0.000000      | 1.000000      | 0.00000      | 0.000000       | 51002.1100   |
| 50%   | 652.000000   | 37.000000    | 5.000000     | 97198.540000  | 1.000000      | 1.00000      | 1.000000       | 100193.9150  |
| 75%   | 718.000000   | 44.000000    | 7.000000     | 127644.240000 | 2.000000      | 1.00000      | 1.000000       | 149388.2470  |
| max   | 850.000000   | 92.000000    | 10.000000    | 250898.090000 | 4.000000      | 1.00000      | 1.000000       | 199992.4800  |

Seems that there might not be any missing data just by looking at counts, however, let's look into missing values anyways in a later section.

But, there likely is outliers if we look at the min and max values of the Balance and EstimatedSalary variables so we will deal with these later.

Now, let's look at the correlation between each variable and see if any pairs are particularly connected with each other.


In [ ]:

```
# Creating a correlation matrix to see what variables are highly/not correlated with each other
correlation_matrix = churn.corr()

# Creating a correlation heatmap to visualize the matrix
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=.5)
plt.title('Correlation Heatmap')
plt.show()
```

<ipython-input-6-c61476b88c26>:2: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
correlation_matrix = churn.corr()
```



There seems to be some sort of negative correlation with Balance and NumOfProducts, Age and Exited, IsActiveMember and Exited.

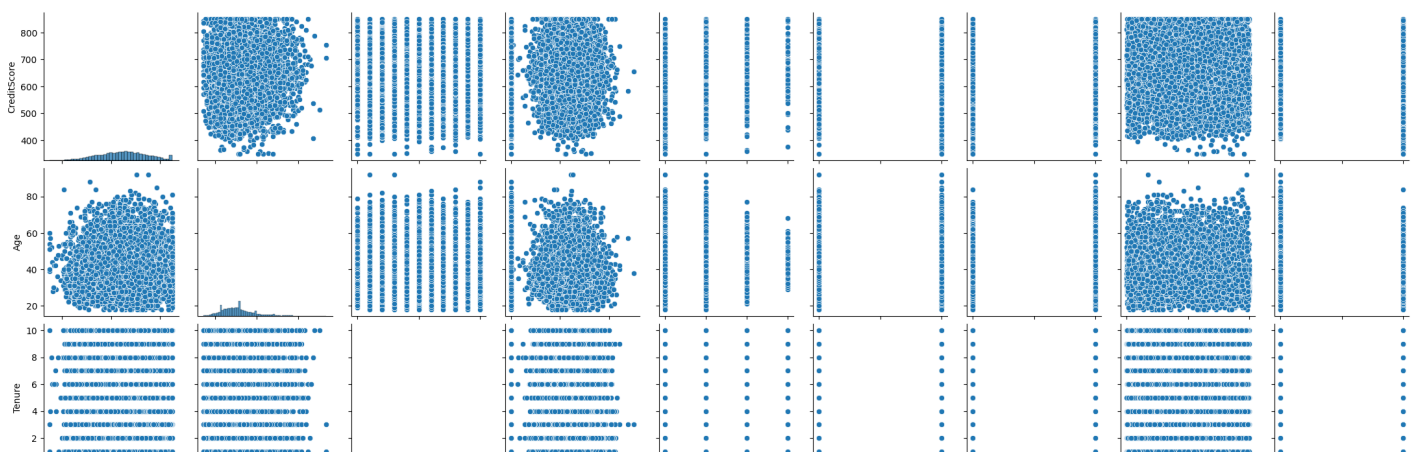
## Visualization

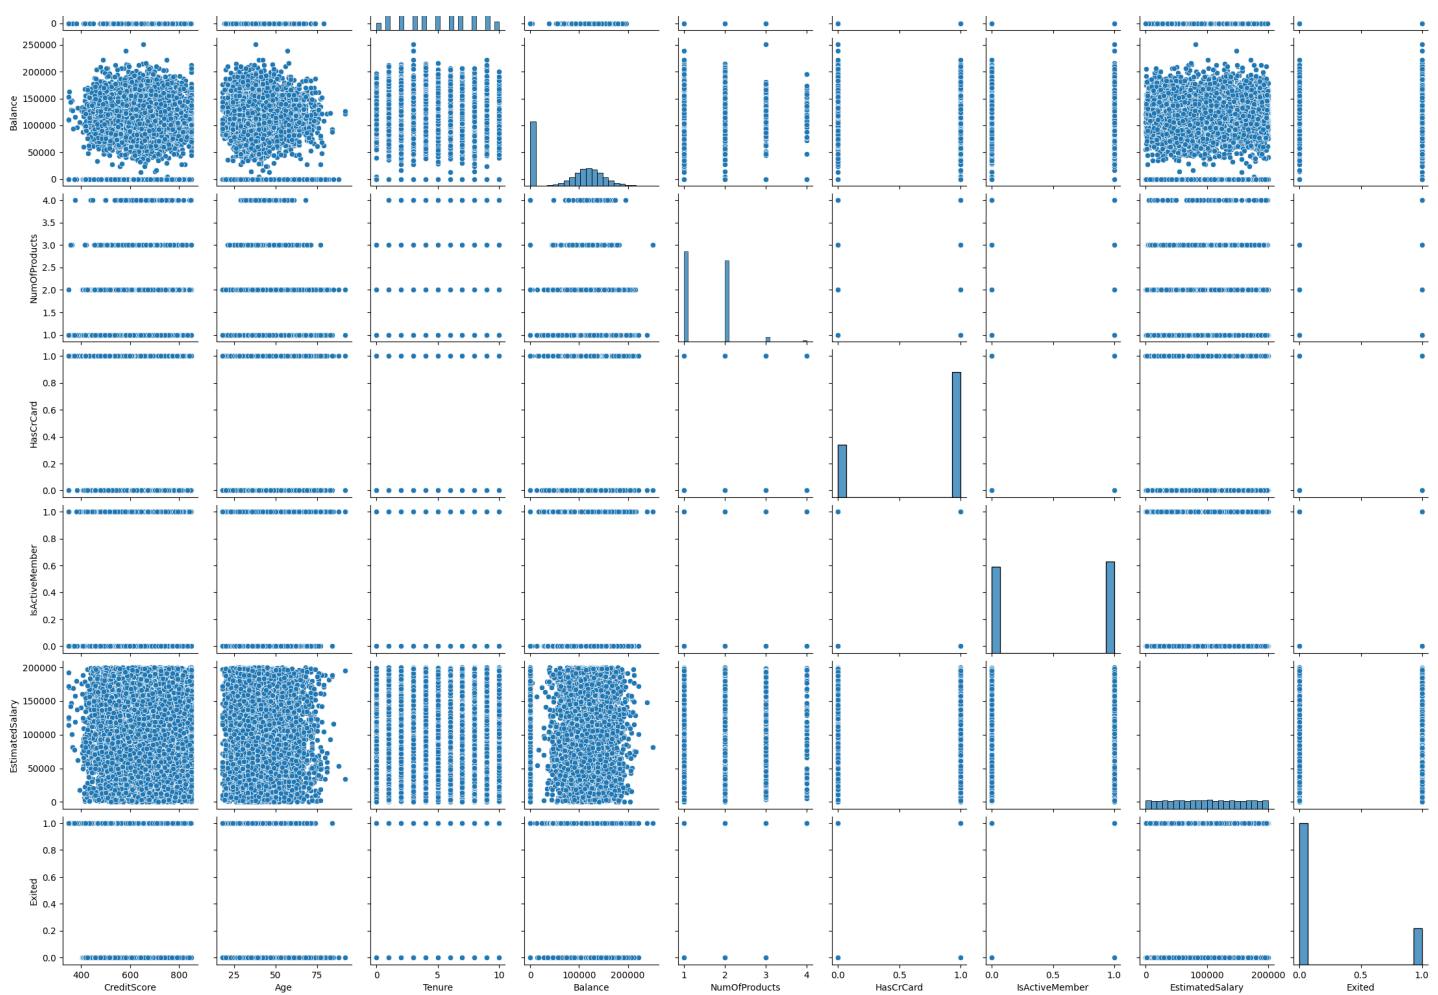
Let's check the distribution of the numerical variables in this dataset

In [ ]:

```
numerical_vars = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
                  'IsActiveMember', 'EstimatedSalary', 'Exited']
numerical_df = churn[numerical_vars]

sns.pairplot(numerical_df) #Pairs each variable with one another to see if there's any specific patterns
plt.show()
```





## Part 2: Data Preprocessing

### Cleaning


Let's see if there are missing values and duplicated rows in our dataset

In [ ]:

```
#Checking for Missing Values
missing_values = churn.isnull().sum()
print("Missing values:\n", missing_values)
ms.matrix(churn)
plt.show()
```

Missing values:

```
CreditScore      0
Geography        0
Gender           0
Age             0
Tenure          0
Balance         0
NumOfProducts   0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64
```



10000

11

11

In [ ]:

```
# Checking duplicated records
duplicates = churn.duplicated().sum()
print('Number of Duplicated Entries: ',duplicates)
```

Number of Duplicated Entries: 0

**Looks like we don't have any null values nor duplicated rows! Cool, we can now move on to the next stage and transform our data as fit**

## Transformation

**The Balance column has some zeroes in each cell which really doesn't make sense, so let's replace them with the mean of the overall column**

In [ ]:


```
# Calculate the mean of non-zero values in the 'Balance' column
balance_mean = churn[churn['Balance'] != 0]['Balance'].mean()


# Replace zeros with the mean value
churn['Balance'] = churn['Balance'].replace(0, balance_mean)
```

In [ ]:

```
# Plot a histogram of the "Balance" column
plt.figure(figsize=(10, 6))
plt.hist(churn['Balance'], bins=30, color='skyblue', edgecolor='black')
plt.title('Distribution of Balance')
plt.xlabel('Balance')
plt.ylabel('Frequency')
plt.grid(True)
plt.show()
```

Distribution of Balance





## Part 3: Model Development and Evaluation

### Model Selection

In [ ]:

```
# Model Libraries
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import SGDClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_curve, auc

# Pipeline Procedures
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
```

In [ ]:

```
X = churn.drop(columns=['Exited']) # setting up our X set that uses every column except
Exited
y = churn['Exited'] # sets up the y set that only contains the Exited variable
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# splits the X and y sets into train/test containing 20% of the dataset randomly
```

### Training and Validation

#### What Metrics are being considered?

- **Accuracy and CV Accuracy:**
  - Accuracy measures the proportion of correct predictions among total predictions
  - Cross-Validation Accuracy is a measure of how well the model performs across multiple subsets of the data using cross-validation
- **Precision:**
  - Measures the proportion of true positive prediction among all positive predictions
- **Recall:**
  - Measures the true positive predictions among all actual positive instances
- **F1-Score:**
  - Harmonic mean of precision and recall
  - Provides a balance between precision and recall and is especially useful when distribution is imbalanced
- **AUC-ROC Curve:**
  - ROC: plot of the true positive rate against the false positive rate
  - AUC: quantifies the overall performance of the model across all possible thresholds

## K-Nearest Neighbor

In [ ]:

```
from sklearn.model_selection import cross_val_score
# Step 1: Define preprocessing steps for numerical and categorical columns
numeric_features = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')), # Handle missing values with median
    ('scaler', StandardScaler()) # Scale features
])

categorical_features = ['Geography', 'Gender']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')), # Handle missing values with a constant value
    ('onehot', OneHotEncoder(handle_unknown='ignore')) # One-hot encode categorical variables
])

# Combine preprocessing steps for numerical and categorical columns
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ])

# Step 2: Create a pipeline with preprocessing and KNN classifier
pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', KNeighborsClassifier())
])

# Step 3: Perform cross-validation
cv_scores = cross_val_score(pipeline, X, y, cv=5) # Perform 5-fold cross-validation

# Step 4: Display cross-validation scores
print("Cross-validation scores:", cv_scores)
print("Mean CV accuracy:", cv_scores.mean())
```

Cross-validation scores: [0.8375 0.84 0.829 0.83 0.829 ]  
Mean CV accuracy: 0.8331

## KNN Model

In [ ]:

```
# Step 1: Define preprocessing steps for numerical and categorical columns
numeric_features = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')), # Handle missing values with median
    ('scaler', StandardScaler()) # Scale features
])

categorical_features = ['Geography', 'Gender']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')), # Handle missing values with a constant value
    ('onehot', OneHotEncoder(handle_unknown='ignore')) # One-hot encode categorical variables
])

# Combine preprocessing steps for numerical and categorical columns
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ])

```



```

# Step 2: Define the KNN classifier
k = 5
knn_classifier = KNeighborsClassifier(n_neighbors=k)

# Step 3: Create the pipeline
pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', knn_classifier)
])

# Step 4: Train the model
pipeline.fit(X_train, y_train)

# Step 5: Evaluate the model
y_pred = pipeline.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Print classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))

```

Accuracy: 0.841

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.88      | 0.93   | 0.90     | 1607    |
| 1            | 0.63      | 0.46   | 0.53     | 393     |
| accuracy     |           |        | 0.84     | 2000    |
| macro avg    | 0.75      | 0.70   | 0.72     | 2000    |
| weighted avg | 0.83      | 0.84   | 0.83     | 2000    |

## ROC Curve

In [ ]:

```

# Step 1: Predict probabilities for the testing data
y_pred_proba = pipeline.predict_proba(X_test)[: , 1]


# Step 2: Compute the False Positive Rate (FPR) and True Positive Rate (TPR) at various thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)


# Step 3: Calculate the Area Under the ROC Curve (AUC-ROC)
auc_roc = auc(fpr, tpr)

# Step 4: Plot the ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (area = {:.2f})'.format(auc_roc))
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()

```

Receiver Operating Characteristic (ROC) Curve





## DecisionTreeClassifier

In [ ]:

```
from sklearn.model_selection import cross_val_score
# Step 1: Define preprocessing steps for numerical and categorical columns
numeric_features = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')), # Handle missing values with median
    ('scaler', StandardScaler()) # Scale features
])

categorical_features = ['Geography', 'Gender']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')), # Handle missing values with a constant value
    ('onehot', OneHotEncoder(handle_unknown='ignore')) # One-hot encode categorical variables
])

# Combine preprocessing steps for numerical and categorical columns
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ])

# Step 2: Create a pipeline with preprocessing and DecisionTreeClassifier
pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', DecisionTreeClassifier())
])

# Step 3: Perform cross-validation
cv_scores = cross_val_score(pipeline, X, y, cv=5) # Perform 5-fold cross-validation

# Step 4: Display cross-validation scores
print("Cross-validation scores:", cv_scores)
print("Mean CV accuracy:", cv_scores.mean())
```

Cross-validation scores: [0.7895 0.786 0.7835 0.781 0.7775]  
Mean CV accuracy: 0.7835

In [ ]:

```
# Step 1: Define preprocessing steps for numerical and categorical columns
numeric_features = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')), # Handle missing values with median
    ('scaler', StandardScaler()) # Scale features
])

categorical_features = ['Geography', 'Gender']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')), # Handle missing values with a constant value
    ('onehot', OneHotEncoder(handle_unknown='ignore')) # One-hot encode categorical variables
])

# Combine preprocessing steps for numerical and categorical columns
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ])

# Step 2: Define the DecisionTreeClassifier
decision_tree_classifier = DecisionTreeClassifier(random_state=42)

# Step 3: Create the pipeline
pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', decision_tree_classifier)
])

# Step 4: Train the model
pipeline.fit(X_train, y_train)

# Step 5: Evaluate the model
y_pred = pipeline.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Print classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.776

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.87      | 0.85   | 0.86     | 1607    |
| 1            | 0.43      | 0.46   | 0.45     | 393     |
| accuracy     |           |        | 0.78     | 2000    |
| macro avg    | 0.65      | 0.66   | 0.65     | 2000    |
| weighted avg | 0.78      | 0.78   | 0.78     | 2000    |


In [ ]:

```
# Step 1: Predict probabilities for the testing data
y_pred_proba = pipeline.predict_proba(X_test)[: , 1]

# Step 2: Compute the False Positive Rate (FPR) and True Positive Rate (TPR) at various thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)

# Step 3: Calculate the Area Under the ROC Curve (AUC-ROC)
auc_roc = auc(fpr, tpr)
```

```
# Step 4: Plot the ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (area = {:.2f})'.format(auc_roc)
)
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```



## Naives Bayes

In [ ]:

```
from sklearn.model_selection import cross_val_score
# Step 1: Define preprocessing steps for numerical and categorical columns
numeric_features = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')), # Handle missing values with median
    ('scaler', StandardScaler()) # Scale features
])

categorical_features = ['Geography', 'Gender']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')), # Handle missing values with a constant value
    ('onehot', OneHotEncoder(handle_unknown='ignore')) # One-hot encode categorical variables
])

# Combine preprocessing steps for numerical and categorical columns
preprocessor = ColumnTransformer(
    transformers=[
```

```

        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ])

# Step 2: Create a pipeline with preprocessing and Naive Bayes classifier
pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', naive_bayes_classifier)
])

# Step 3: Perform cross-validation
cv_scores = cross_val_score(pipeline, X, y, cv=5) # Perform 5-fold cross-validation

# Step 4: Display cross-validation scores
print("Cross-validation scores:", cv_scores)
print("Mean CV accuracy:", cv_scores.mean())

Cross-validation scores: [0.797  0.801  0.808  0.807  0.7955]
Mean CV accuracy: 0.8017

```

In [ ]:

```

# Step 1: Define preprocessing steps for numerical and categorical columns
numeric_features = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')), # Handle missing values with median
    ('scaler', StandardScaler()) # Scale features
])

categorical_features = ['Geography', 'Gender']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')), # Handle missing values with a constant value
    ('onehot', OneHotEncoder(handle_unknown='ignore')) # One-hot encode categorical variables
])

# Combine preprocessing steps for numerical and categorical columns
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ])

# Step 2: Define the Gaussian Naive Bayes classifier
naive_bayes_classifier = GaussianNB()

# Step 3: Create the pipeline
pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', naive_bayes_classifier)
])

# Step 4: Train the model
pipeline.fit(X_train, y_train)

# Step 5: Evaluate the model
y_pred = pipeline.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Print classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))

```

```

Accuracy: 0.8125
Classification Report:
              precision    recall  f1-score   support

```

|              |   |      |      |      |      |
|--------------|---|------|------|------|------|
|              | 0 | 0.85 | 0.93 | 0.89 | 1607 |
|              | 1 | 0.54 | 0.35 | 0.42 | 393  |
| accuracy     |   |      |      | 0.81 | 2000 |
| macro avg    |   | 0.69 | 0.64 | 0.66 | 2000 |
| weighted avg |   | 0.79 | 0.81 | 0.80 | 2000 |


In [ ]:

```
# Step 1: Predict probabilities for the testing data
y_pred_proba = pipeline.predict_proba(X_test)[: , 1]

# Step 2: Compute the False Positive Rate (FPR) and True Positive Rate (TPR) at various thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)

# Step 3: Calculate the Area Under the ROC Curve (AUC-ROC)
auc_roc = auc(fpr, tpr)

# Step 4: Plot the ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (area = {:.2f})'.format(auc_roc))
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```



## RandomForestClassifier

In [ ]:

```

from sklearn.model_selection import cross_val_score
# Step 1: Define preprocessing steps for numerical and categorical columns
numeric_features = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')), # Handle missing values with median
    ('scaler', StandardScaler()) # Scale features
])

categorical_features = ['Geography', 'Gender']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')), # Handle missing values with a constant value
    ('onehot', OneHotEncoder(handle_unknown='ignore')) # One-hot encode categorical variables
])

# Combine preprocessing steps for numerical and categorical columns
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ])

# Step 2: Create a pipeline with preprocessing and RandomForest classifier
pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', random_forest_classifier)])

# Step 3: Perform cross-validation
cv_scores = cross_val_score(pipeline, X, y, cv=5) # Perform 5-fold cross-validation

# Step 4: Display cross-validation scores
print("Cross-validation scores:", cv_scores)
print("Mean CV accuracy:", cv_scores.mean())

```

Cross-validation scores: [0.8475 0.848 0.8485 0.8465 0.855 ]  
Mean CV accuracy: 0.8491

In [ ]:

```

# Step 1: Define preprocessing steps for numerical and categorical columns
numeric_features = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')), # Handle missing values with median
    ('scaler', StandardScaler()) # Scale features
])

categorical_features = ['Geography', 'Gender']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')), # Handle missing values with a constant value
    ('onehot', OneHotEncoder(handle_unknown='ignore')) # One-hot encode categorical variables
])

# Combine preprocessing steps for numerical and categorical columns
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ])

# Step 2: Define the RandomForestClassifier
random_forest_classifier = RandomForestClassifier(random_state=42)

# Step 3: Create the pipeline
pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', random_forest_classifier)
])

```

```

# Step 4: Train the model
pipeline.fit(X_train, y_train)

# Step 5: Evaluate the model
y_pred = pipeline.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Print classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))

```

Accuracy: 0.8525

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.88      | 0.95   | 0.91     | 1607    |
| 1            | 0.69      | 0.45   | 0.54     | 393     |
| accuracy     |           |        | 0.85     | 2000    |
| macro avg    | 0.78      | 0.70   | 0.73     | 2000    |
| weighted avg | 0.84      | 0.85   | 0.84     | 2000    |

In [ ]:

```


# Step 1: Predict probabilities for the testing data
y_pred_proba = pipeline.predict_proba(X_test)[: , 1]

# Step 2: Compute the False Positive Rate (FPR) and True Positive Rate (TPR) at various thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)


# Step 3: Calculate the Area Under the ROC Curve (AUC-ROC)
auc_roc = auc(fpr, tpr)

# Step 4: Plot the ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (area = {:.2f})'.format(auc_roc))
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()

```







## Stochastic Gradient Descent

In [ ]:

```
from sklearn.model_selection import cross_val_score
# Step 1: Define preprocessing steps for numerical and categorical columns
numeric_features = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')), # Handle missing values with median
    ('scaler', StandardScaler()) # Scale features
])

categorical_features = ['Geography', 'Gender']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')), # Handle missing values with a constant value
    ('onehot', OneHotEncoder(handle_unknown='ignore')) # One-hot encode categorical variables
])

# Combine preprocessing steps for numerical and categorical columns
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ])

# Step 2: Create a pipeline with preprocessing and KNN classifier
pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', SGDClassifier())
])

# Step 3: Perform cross-validation
cv_scores = cross_val_score(pipeline, X, y, cv=5) # Perform 5-fold cross-validation

# Step 4: Display cross-validation scores
print("Cross-validation scores:", cv_scores)
print("Mean CV accuracy:", cv_scores.mean())
```

Cross-validation scores: [0.796 0.796 0.7965 0.791 0.7905]  
Mean CV accuracy: 0.7939999999999999

In [ ]:

```
# Step 1: Define preprocessing steps for numerical and categorical columns
numeric_features = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')), # Handle missing values with median
    ('scaler', StandardScaler()) # Scale features
])

categorical_features = ['Geography', 'Gender']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')), # Handle missing values with a constant value
    ('onehot', OneHotEncoder(handle_unknown='ignore')) # One-hot encode categorical variables
])

# Combine preprocessing steps for numerical and categorical columns
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ])

# Step 2: Create a pipeline with preprocessing and KNN classifier
pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', SGDClassifier())
])

# Step 3: Perform cross-validation
cv_scores = cross_val_score(pipeline, X, y, cv=5) # Perform 5-fold cross-validation

# Step 4: Display cross-validation scores
print("Cross-validation scores:", cv_scores)
print("Mean CV accuracy:", cv_scores.mean())
```

```

ing values with a constant value
    ('onehot', OneHotEncoder(handle_unknown='ignore')) # One-hot encode categorical variables
])

# Combine preprocessing steps for numerical and categorical columns
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ])

# Step 2: Define the SGDClassifier
sgd_classifier = SGDClassifier(random_state=42)

# Step 3: Create the pipeline
pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', sgd_classifier)
])

# Step 4: Train the model
pipeline.fit(X_train, y_train)

# Step 5: Evaluate the model
y_pred = pipeline.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Print classification report
print("Classification Report:")
print(classification_report(y_test, y_pred, zero_division=0))

```

Accuracy: 0.8035

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.80      | 1.00   | 0.89     | 1607    |
| 1            | 0.00      | 0.00   | 0.00     | 393     |
| accuracy     |           |        | 0.80     | 2000    |
| macro avg    | 0.40      | 0.50   | 0.45     | 2000    |
| weighted avg | 0.65      | 0.80   | 0.72     | 2000    |

In [ ]:

```

# Step 1: Predict probabilities for the testing data
y_pred_proba = pipeline.predict_proba(X_test)[: , 1]


# Step 2: Compute the False Positive Rate (FPR) and True Positive Rate (TPR) at various thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)

# Step 3: Calculate the Area Under the ROC Curve (AUC-ROC)
auc_roc = auc(fpr, tpr)

# Step 4: Plot the ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (area = {:.2f})'.format(auc_roc))
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()

```

Receiver Operating Characteristic (ROC) Curve



## LogisticRegression

In [ ]:

```
from sklearn.model_selection import cross_val_score
# Step 1: Define preprocessing steps for numerical and categorical columns
numeric_features = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')), # Handle missing values with median
    ('scaler', StandardScaler()) # Scale features
])

categorical_features = ['Geography', 'Gender']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')), # Handle missing values with a constant value
    ('onehot', OneHotEncoder(handle_unknown='ignore')) # One-hot encode categorical variables
])

# Combine preprocessing steps for numerical and categorical columns
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ]
)

# Step 2: Create a pipeline with preprocessing and KNN classifier
pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', logistic_regression)
])

# Step 3: Perform cross-validation
cv_scores = cross_val_score(pipeline, X, y, cv=5) # Perform 5-fold cross-validation
```

```
# Step 4: Display cross-validation scores
print("Cross-validation scores:", cv_scores)
print("Mean CV accuracy:", cv_scores.mean())
```

```
Cross-validation scores: [0.792  0.792  0.793  0.7945 0.798 ]
Mean CV accuracy: 0.7939
```

In [ ]:

```
# Step 1: Define preprocessing steps for numerical and categorical columns
numeric_features = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')), # Handle missing values with median
    ('scaler', StandardScaler()) # Scale features
])

categorical_features = ['Geography', 'Gender']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')), # Handle missing values with a constant value
    ('onehot', OneHotEncoder(handle_unknown='ignore')) # One-hot encode categorical variables
])

# Combine preprocessing steps for numerical and categorical columns
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ])

# Step 2: Define the Logistic Regression model
logistic_regression = LogisticRegression(random_state=42)

# Step 3: Create the pipeline
pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', logistic_regression)
])

# Step 4: Train the model
pipeline.fit(X_train, y_train)

# Step 5: Evaluate the model
y_pred = pipeline.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Print classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

```
Accuracy: 0.8045
```

```
Classification Report:
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.82      | 0.97   | 0.89     | 1607    |
| 1            | 0.51      | 0.14   | 0.22     | 393     |
| accuracy     |           |        | 0.80     | 2000    |
| macro avg    | 0.67      | 0.55   | 0.56     | 2000    |
| weighted avg | 0.76      | 0.80   | 0.76     | 2000    |

In [ ]:

```
# Step 1: Predict probabilities for the testing data
y_pred_proba = pipeline.predict_proba(X_test)[:, 1]
```


```

# Step 2: Compute the False Positive Rate (FPR) and True Positive Rate (TPR) at various thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)

# Step 3: Calculate the Area Under the ROC Curve (AUC-ROC)
auc_roc = auc(fpr, tpr)

# Step 4: Plot the ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (area = {:.2f})'.format(auc_roc))
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()

```



## Which Model Is Best?

### Cross-Validation Scores

- **Best Mean Cross-Validation Score: RandomForest**

### Metric Performance

- **Accuracy Score**
  - **Highest Accuracy Score: RandomForest**
- **Precision**
  - **Highest Weighted Average: RandomForest**
- **Recall**
  - **Highest Weighted Average: RandomForest**
- **F1-Score**

- **Highest Weighted Average: RandomForest**

#### **ROC Curve Areas**

- **Best Area Under Curve: RandomForest**

**Overall, it seems that the best model to use for the Churn dataset would be the RandomForestClassifier model**